

**Your grade: 100%**[Next item →](#)Your latest: **92.50%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

1. What do you think applying this filter to a grayscale image will do?

**1 / 1 point**

$$\begin{bmatrix} 0 & 1 & 1 & 0 \\ 1 & 3 & 3 & 1 \\ -1 & -3 & -3 & -1 \\ 0 & -1 & -1 & 0 \end{bmatrix}$$

- ☒ Detect horizontal edges.
- ☐ Detecting image contrast.
- ☐ Detect 45-degree edges.
- ☐ Detect vertical edges.

 **Correct**

Correct. There is a high difference between the values in the top part from those in the bottom part of the matrix. When convolving this filter on a grayscale image, the horizontal edges will be detected.

2. Suppose your input is a 128 by 128 grayscale image, and you are not using a convolutional network. If the first hidden layer has 256 neurons, each one fully connected to the input, how many parameters does this hidden layer have (including the bias parameters)?

**1 / 1 point**

- ☒ 4194560
- ☐ 4194304
- ☐ 12582912



☐ 12583168

☒ **Correct**

Correct, the number of inputs for each unit is  $128 \times 128$  since the input image is grayscale, so we need  $128 \times 128 \times 256$  parameters for the weights and 256 parameters for the bias thus  $128 \times 128 \times 256 + 256 = 4194560$ .

3. Suppose your input is a 256 by 256 color (RGB) image, and you use a convolutional layer with 128 filters that are each  $7 \times 7$ . How many parameters does this hidden layer have (including the bias parameters)?

1 / 1 point

☐ 6400

☒ 18944

☐ 1233125504

☐ 18816

☒ **Correct**

Yes, you have  $7 \times 7 \times 3 + 1$  weights per filter with the bias. Given that you have 128 filters, you get  $(7 \times 7 \times 3 + 1) \times 128 = 18944$ .

4. You have an input volume that is  $121 \times 121 \times 16$ , and convolve it with 32 filters of  $4 \times 4$ , using a stride of 3 and no padding. What is the output volume?

1 / 1 point

☒  $40 \times 40 \times 32$

☐  $40 \times 40 \times 16$

☐  $118 \times 118 \times 16$

☐  $118 \times 118 \times 32$



✓ **Correct**

Correct, using the formula  $n_H^{[l]} = \frac{n_H^{[l-1]} + 2 \times p - f}{s} + 1$  with  $n_H^{[l-1]} = 121$ ,  $p = 0$ ,  $f = 4$ , and  $s = 3$  we get 40

5. You have an input volume that is 61x61x32, and pad it using “pad=3”. What is the dimension of the resulting volume (after padding)?

1 / 1 point

- ☒ 67x67x32
- ☐ 64x64x35
- ☐ 64x64x32
- ☐ 61x61x35

✓ **Correct**

Yes, if the padding is 3 you add 6 to the height dimension and 6 to the width dimension.

6. You have an input volume that is 63x63x16, and convolve it with 32 filters that are each 7x7, and stride of 1. You want to use a “same” convolution. What is the padding?

1 / 1 point

- ☒ 3
- ☐ 2
- ☐ 1
- ☐ 7

✓ **Correct**

Correct, you need to satisfy the following equation:  $n_H - f + 2 \times p + 1 = n_H$  as you want to



keep the dimensions between the input volume and the output volume.

7. You have an input volume that is 32x32x16, and apply max pooling with a stride of 2 and a filter size of 2. What is the output volume?

1 / 1 point

- ☐ 16x16x8
- ☒ 16x16x16
- ☐ 32x32x8
- ☐ 15x15x16

✓ **Correct**

Correct, using the following formula:  $n_H^{[l]} = \frac{n_H^{[l-1]} + 2 \times p - f}{s} + 1$

8. Which of the following are hyperparameters of the pooling layers? (Choose all that apply)

0.8 / 1 point

☒ Whether it is max or average.

✓ **Correct**

Yes, these are the two types of pooling discussed in the lectures, and choosing which to use is considered a hyperparameter.

☒ Filter size.

✓ **Correct**

Yes, although usually, we set  $f = s$  this is one of the hyperparameters of a pooling layer.

☐ Average weights.



☒ Number of filters.

 **This should not be selected**

No, pooling layers keep the depth dimension of the volume, we don't need to specify the number of filters with a pooling layer.

9. In lecture we talked about “parameter sharing” as a benefit of using convolutional networks. Which of the following statements about parameter sharing in ConvNets are true? (Check all that apply)

0.5 / 1 point

☒ It allows parameters learned for one task to be shared even for a different task (transfer learning).

 **This should not be selected**

No, transfer learning is not bound to ConvNets and can be used with other types of models as you've seen in Course 1-3.

☐ It reduces the total number of parameters, thus reducing overfitting.

☐ It allows gradient descent to set many of the parameters to zero, thus making the connections sparse.

☒ It allows a feature detector to be used in multiple locations throughout the whole input image/input volume.

 **Correct**

Yes, by sliding a filter of parameters over the entire input volume, we make sure a feature detector can be used in multiple locations.

10. In lecture we talked about “sparsity of connections” as a benefit of using convolutional layers. What does this mean?

1 / 1 point

☐ Regularization causes gradient descent to set many of the parameters to zero.



✗ **This should not be selected**

No, transfer learning is not bound to ConvNets and can be used with other types of models as you've seen in Course 1-3.

- ☐ It reduces the total number of parameters, thus reducing overfitting.
- ☐ It allows gradient descent to set many of the parameters to zero, thus making the connections sparse.
- ☒ It allows a feature detector to be used in multiple locations throughout the whole input image/input volume.

✓ **Correct**

Yes, by sliding a filter of parameters over the entire input volume, we make sure a feature detector can be used in multiple locations.

10. In lecture we talked about “sparsity of connections” as a benefit of using convolutional layers. What does this mean?

1 / 1 point

- ☐ Regularization causes gradient descent to set many of the parameters to zero.
- ☐ Each layer in a convolutional network is connected only to two other layers
- ☐ Each filter is connected to every channel in the previous layer.
- ☒ Each activation in the next layer depends on only a small number of activations from the previous layer.

✓ **Correct**

Yes, each activation of the output volume is computed by multiplying the parameters from **only one filter** with a volumic slice of the input volume and then summing all these together.